



Scan for the package,
the do-file, and this poster

Coding Open-Ended Survey Text Without Leaving Stata: the catllm Package for LLM-Assisted Content Coding

Chris Soria¹

¹Department of Demography, University of California, Berkeley



Abstract


Open-ended survey items capture what closed-ended items cannot, and are routinely dropped from analysis because hand-coding them does not scale. `catllm` is a Stata command that codes free-text responses with large language models and returns ordinary Stata indicator variables. It wraps a Python backend (`cat-stack`) behind five verbs: `extract`/`explore` find categories, `classify` applies them, and `summarize`/`setup` close the loop, across OpenAI, Anthropic, Google, or Ollama. Tested on multi-label, multi-class items, proprietary models agree with human coders on **97%** of straightforward items and **88–91%** of complex interpretive ones; a multi-model ensemble reaches **98%** agreement with human consensus.

The Problem

- Researchers already use LLMs to code open-ended text, but little guidance exists for keeping classifications aligned with intended meaning without skewing conclusions.
- Model choice is often arbitrary. Researchers overpay for models that excel at general tasks, which rarely predicts coding accuracy.
- Coding several open-ended items at once means devising a bespoke strategy each time. That risks bias and error, not a documented method.
- Doing this from Stata isn't supported today: researchers export to CSV, script in Python, re-merge, and lose the audit trail that makes coding reproducible.
- Aim:** put the whole loop inside Stata, and tackle over-classification from several evidence-based angles: category definitions, model choice, and cross-provider ensemble consensus.

The Pipeline

Text goes in as a variable — categories come back as variables



Categories come from `extract` or your own list; `classify` returns multi-label indicator variables either way.

`generate()` is a prefix, not a variable name: a five-category scheme returns five 0/1 byte variables. `extract`'s specificity, iteration count, and category limit are tunable, so discovery matches the research question instead of a fixed default.

One Command

```
local cats ///
"Housing cost: Moved over rent, a mortgage, or affordability." ///
"Employment: Moved for a job, a transfer, or school." ///
"Family: Moved for caregiving or to be near relatives." ///
"Other: Does not fit any of the above categories."
```

```
catllm classify response, ///
categories(`cats') ///
apikey($OPENAI_API_KEY) ///
model("gpt-4o-mini") ///
generate(reason)
```


```
tab1 reason_*
regress reason_Housing_cost age female
```

Category definitions, not one-word labels: writing each category as a sentence classifies several percentage points better on every model tested. An explicit `other` category is required, not optional — without it the model must guess on responses that genuinely fit nothing.

References

Soria (2026). *CatLLM: A Python package for generating, assigning, and scoring open-ended survey data and images*. | Soria (2026). *High agreement, different stories*. | Soria (2026). *Model diversity over model size*. | Bojic et al. (2025). *Comparing LLM and human annotators*. | Wei et al. (2023). *Chain-of-thought prompting*. | Dhuliawala et al. (2023). *Chain-of-verification*. | Zheng et al. (2024). *Step-back prompting*.

Accuracy Against Human Coders




UC Berkeley Social Networks Study validation: macro F1 against human consensus, averaged across three survey questions so no single question dominates. Frontier-closed, economy-closed, and frontier-open tiers overlap closely; Claude Fable 5, shown individually, lands in line with the top cloud tiers. Ranges span the models within each tier.

What Moves Accuracy

- Cloud tier spread is narrow.** Frontier-closed, economy-closed, and frontier-open models span just 0.07 of macro F1 (0.75–0.83); Claude Fable 5 lands at 0.81, in line with the top tiers.
- Response brevity helps.** Responses under 50 characters are classified 7–11 points more accurately than long ones.
- Demographic variation is minimal.** Accuracy is stable across population segments — an essential precondition for using coded output in stratified analysis.
- Structured output is not the bottleneck.** Across 8 models × 3,208 responses (25,664 classifications), 100% of successful calls returned valid structured output; the 0–18 failures per model were transient server errors, not parsing failures.

Local Model Spread




Same macro F1 against human consensus as the cloud tiers, for the four laptop-scale Ollama models. The spread here (0.62–0.77) is more than double the cloud tiers' (0.75–0.83) — picking well matters far more once you're local.

Choosing a Route

Route	Cost	Speed	Data leaves machine
Frontier proprietary	High	Slow	Yes
Small proprietary	Low	Fast	Yes
Multi-vendor ensemble	Sum of parts	Slowest member	Yes
Local open-weight (Ollama)	\$0	Hardware-bound	No

For restricted-use data under an IRB or data-use agreement, the local route is often the only admissible one — and it is the route the package defaults to protecting, auto-starting the local server and enabling JSON recovery for smaller models.

Ensemble & Consensus Voting



`models()` takes several model/provider pairs; `consensus()` sets the bar for assigning a category. Cloud and local models can be mixed freely in one ensemble. Alongside each label the package returns an agreement score, so the rows where models split are identifiable rather than silently averaged away.

98%

ensemble vs.
human consensus

25,664

benchmark
classifications

\$0

fully local
route

Key Findings

• 0.75–0.83

Macro F1 against human consensus across cloud tiers; vendor and tier matter far less than the local-vs-cloud choice covered in Choosing a Route

• 98%

A three-model ensemble agrees with human consensus more often than any single model in it — disagreement between models is informative, not noise

• 0.62–0.77

Macro F1 across the four local Ollama models — more than double the cloud tiers' spread, so which local model you pick matters a lot

Reproduce This

Everything on this poster runs from one self-contained do-file — no data downloads, no file paths, guarded so it completes with a single API key or with none at all:

```
net install catllm, ///
from("https://raw.githubusercontent.com/chrissoria/cat-llm/main/stata-package/") ///
replace

catllm setup
do catllm_stata_conference.do
```

The file walks the full pipeline: classify, analyze the resulting indicators, compare three providers, run an ensemble with agreement scores, discover a coding scheme with `extract`, and repeat the whole thing locally at zero cost. Eight further worked examples ship with the package.

Limitations

- Concordance with human coders is not validity.** Where humans systematically misread a response, an agreeing model does too; agreement measures replication of a benchmark, not correctness.
- Validation is strongest on English-language survey text** of the kind studied here. Performance on other languages, registers, and domains needs separate checking on your own data.
- Models are moving targets.** A model version can be deprecated or silently updated, so the version string belongs in your methods section alongside the seed.
- Smaller local models are less tested.** The published benchmarks use high-end models; the low end warrants its own validation before use in production coding.